

Verifying Parallel Programs with MPI-SPIN

Part 4: Numerical Computation

Stephen F. Siegel

Department of Computer and Information Sciences
University of Delaware

EuroPVM/MPI 2007
Paris, France
30 September 2007

Overview

1. Goal: prove that program computes correct result
2. Symbolic execution
 - Performing symbolic arithmetic in MPI-SPIN
3. Functional equivalence
4. Three types of numerical equivalence
5. Diffusion revisited
6. Dealing with branches: the path condition

Goal: prove that program computes correct result

- what does it mean to say program computes correct result?

Goal: prove that program computes correct result

- what does it mean to say program computes correct result?
- verification requires **specification**
 - definition of *correct*

Goal: prove that program computes correct result

- what does it mean to say program computes correct result?
- verification requires **specification**
 - definition of *correct*
- our specification
 - a trusted **sequential** version of program

Goal: prove that program computes correct result

- what does it mean to say program computes correct result?
- verification requires **specification**
 - definition of *correct*
- our specification
 - a trusted **sequential** version of program
- our method
 - use MPI-SPIN to prove the sequential and parallel program are **functionally equivalent**
 - i.e., produce same output for any given input
 - reduces the problem of verifying the correctness of a **parallel** numerical program to the problem of verifying the correctness of a **sequential** numerical program
 - uses symbolic execution to model floating-point computation

How do we model floating-point computation?

- one double-precision floating-point variable has 2^{64} possible states
- abstraction?

Performing symbolic arithmetic in MPI-SPIN

- type
 - `MPI_Symbolic`

Performing symbolic arithmetic in MPI-SPIN

- type
 - `MPI_Symbolic`
- constants of type `MPI_Symbolic`
 1. `SYM_ZERO`: zero (0)
 2. `SYM_ONE`: one (1)
 3. `SYM_FALSE`: the boolean value `false`
 4. `SYM_TRUE`: the boolean value `true`

Forming new symbolic expressions from old

- the following return an `MPI_Symbolic` of **numeric** type
 - `SYM_add(MPI_Symbolic x, MPI_Symbolic y)`
 - `SYM_subtract(MPI_Symbolic x, MPI_Symbolic y)`
 - `SYM_multiply(MPI_Symbolic x, MPI_Symbolic y)`
 - `SYM_divide(MPI_Symbolic x, MPI_Symbolic y)`
 - `SYM_sqrt(MPI_Symbolic x)`
 - `SYM_abs(MPI_Symbolic x)`
 - `SYM_if(MPI_Symbolic b, MPI_Symbolic x, MPI_Symbolic y)`
 - `(b ? x : y)`

How does MPI-SPIN represent symbolic expressions?

Problem

- symbolic expressions can get big
- there can be millions of states
- storing all symbolic expressions in every state would quickly consume all memory

How does MPI-SPIN represent symbolic expressions?

Problem

- symbolic expressions can get big
- there can be millions of states
- storing all symbolic expressions in every state would quickly consume all memory

Solution: Value numbering

- place all symbolic expressions in a shared **expression table**
 - every expression has a unique **ID number**

How does MPI-SPIN represent symbolic expressions?

Problem

- symbolic expressions can get big
- there can be millions of states
- storing all symbolic expressions in every state would quickly consume all memory

Solution: Value numbering

- place all symbolic expressions in a shared **expression table**
 - every expression has a unique **ID number**
- in the model...
 - replace all floating-point values with ID numbers

How does MPI-SPIN represent symbolic expressions?

Problem

- symbolic expressions can get big
- there can be millions of states
- storing all symbolic expressions in every state would quickly consume all memory

Solution: Value numbering

- place all symbolic expressions in a shared **expression table**
 - every expression has a unique **ID number**
- in the model...
 - replace all floating-point values with ID numbers
 - replace all floating-point operations with symbolic operations
 - to evaluate $x + y$:
 - is $x + y$ already in the table?
 - if yes, return its ID number
 - if no, create new table entry and return new ID number

Symbolic expression table: example

i	e_i	interpretation
-----	-------	----------------

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$
21	(+, 19, 20)	$(0.0 + x_2 x_4) + x_3 x_6$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$
21	(+, 19, 20)	$(0.0 + x_2 x_4) + x_3 x_6$
22	(*, 4, 7)	$x_2 x_5$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$
21	(+, 19, 20)	$(0.0 + x_2 x_4) + x_3 x_6$
22	(*, 4, 7)	$x_2 x_5$
23	(+, 0, 22)	$0.0 + x_2 x_5$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$
21	(+, 19, 20)	$(0.0 + x_2 x_4) + x_3 x_6$
22	(*, 4, 7)	$x_2 x_5$
23	(+, 0, 22)	$0.0 + x_2 x_5$
24	(*, 5, 9)	$x_3 x_7$

Symbolic expression table: example

i	e_i	interpretation
0	(L, 0.0)	0.0
1	(L, 1.0)	1.0
2	(X, 0)	x_0
3	(X, 1)	x_1
4	(X, 2)	x_2
5	(X, 3)	x_3
6	(X, 4)	x_4
7	(X, 5)	x_5
8	(X, 6)	x_6
9	(X, 7)	x_7
10	(*, 2, 6)	$x_0 x_4$
11	(+, 0, 10)	$0.0 + x_0 x_4$
12	(*, 3, 8)	$x_1 x_6$

i	e_i	interpretation
13	(+, 11, 12)	$(0.0 + x_0 x_4) + x_1 x_6$
14	(*, 2, 7)	$x_0 x_5$
15	(+, 0, 14)	$0.0 + x_0 x_5$
16	(*, 3, 9)	$x_1 x_7$
17	(+, 15, 16)	$(0.0 + x_0 x_5) + x_1 x_7$
18	(*, 4, 6)	$x_2 x_4$
19	(+, 0, 12)	$0.0 + x_2 x_4$
20	(*, 5, 8)	$x_3 x_6$
21	(+, 19, 20)	$(0.0 + x_2 x_4) + x_3 x_6$
22	(*, 4, 7)	$x_2 x_5$
23	(+, 0, 22)	$0.0 + x_2 x_5$
24	(*, 5, 9)	$x_3 x_7$
25	(+, 23, 24)	$(0.0 + x_2 x_5) + x_3 x_7$

Functional equivalence

- Goal
 - prove sequential and parallel programs are functionally equivalent

Functional equivalence

- Goal
 - prove sequential and parallel programs are functionally equivalent
- Method
 1. construct symbolic model M_{seq} of sequential program
 - input: $\mathbf{x} = (x_1, \dots, x_n)$, output: \mathbf{y}

Numerical Issues

Problem: distinct symbolic expressions should be considered “**equivalent**” in some cases

Example: **real equivalence**

- $((x_3 + x_1) + x_2) + x_0$ and $((x_0 + x_1) + x_2) + x_3$

Numerical Issues

Problem: distinct symbolic expressions should be considered “**equivalent**” in some cases

Example: **real equivalence**

- $((x_3 + x_1) + x_2) + x_0$ and $((x_0 + x_1) + x_2) + x_3$
- if computer arithmetic were real arithmetic then evaluating these expressions would yield identical results for any values of x_0, x_1, \dots

Numerical Issues

Problem: distinct symbolic expressions should be considered “**equivalent**” in some cases

Example: **real equivalence**

- $((x_3 + x_1) + x_2) + x_0$ and $((x_0 + x_1) + x_2) + x_3$
- if computer arithmetic were real arithmetic then evaluating these expressions would yield identical results for any values of x_0, x_1, \dots
- computer arithmetic is not real arithmetic
 - e.g., floating-point addition can never be associative
 - **What every computer scientist should know about floating-point arithmetic**
 - David Goldberg
 - ACM Computing Surveys 23(1), 1991

Numerical Issues, cont.

- in some cases, knowing the sequential and parallel programs are “**real equivalent**” is good enough
- if testing yields slightly different results...
 - ...but you know programs are real equivalent
 - then you know the **only reason** results differ is due to vagaries of floating-point arithmetic
 - and **not** to some error in your parallel program

Three equivalence relations

- MPI-SPIN supports three different **equivalence relations** on the set of symbolic expressions
 0. Herbrand
 1. IEEE
 2. Real

Three equivalence relations

- MPI-SPIN supports three different **equivalence relations** on the set of symbolic expressions
 0. Herbrand
 1. IEEE
 2. Real
- user specifies which numeric mode to use at command-line
 - `ms -sym=0 ...` for Herbrand, etc.

Three equivalence relations

- MPI-SPIN supports three different **equivalence relations** on the set of symbolic expressions
 0. Herbrand
 1. IEEE
 2. Real
- user specifies which numeric mode to use at command-line
 - `ms -sym=0 ...` for Herbrand, etc.
- when a new expression is formed it is reduced to a **canonical form** before being inserted into table
 - goal is for each equivalence class to have at most one representative in table

Herbrand equivalence

- two symbolic expressions are **Herbrand equivalent** iff they are identical

Herbrand equivalence

- two symbolic expressions are **Herbrand equivalent** iff they are identical
- numeric operations are treated as **uninterpreted functions**
- example: x and $x + 0$ are not Herbrand equivalent

Herbrand equivalence

- two symbolic expressions are **Herbrand equivalent** iff they are identical
- numeric operations are treated as **uninterpreted functions**
- example: x and $x + 0$ are not Herbrand equivalent
- Herbrand equivalence is the **strongest** form of equivalence
 - if two programs are Herbrand equivalent then they will produce the same result no matter how numeric operations are implemented
 - as long as the numeric operations are deterministic functions!
- in many complex examples, sequential and parallel versions are Herbrand equivalent
 - complexity lies elsewhere (e.g., in distribution of data, coordination of processes)

IEEE equivalence

- two expressions are **IEEE equivalent** if one can be reduced to the other using the following identities:
 - $x + y = y + x$
 - $x + 0 = x = 0 + x$
 - $x - x = 0$
 - $xy = yx$
 - $1x = x = x1$
 - $x/x = 1$ (if $x \neq 0$)
 - \vdots

IEEE equivalence

- two expressions are **IEEE equivalent** if one can be reduced to the other using the following identities:
 - $x + y = y + x$
 - $x + 0 = x = 0 + x$
 - $x - x = 0$
 - $xy = yx$
 - $1x = x = x1$
 - $x/x = 1$ (if $x \neq 0$)
 - \vdots
- rationale
 - all of these identities are guaranteed to hold on any platform conforming to the IEEE-754 standard

IEEE equivalence

- two expressions are **IEEE equivalent** if one can be reduced to the other using the following identities:
 - $x + y = y + x$
 - $x + 0 = x = 0 + x$
 - $x - x = 0$
 - $xy = yx$
 - $1x = x = x1$
 - $x/x = 1$ (if $x \neq 0$)
 - \vdots
- rationale
 - all of these identities are guaranteed to hold on any platform conforming to the IEEE-754 standard
- IEEE equivalence is **weaker** than Herbrand equivalence
- two IEEE equivalent programs are guaranteed to produce the exact same results if executed on an IEEE-754-compliant platform
 - but results may differ on non-compliant platforms

Real equivalence

- two expressions are **real equivalent** if one can be reduced to the other using any of the field identities
 - all of the IEEE identities
 - associativity of addition and multiplication
 - $x(1/x) = 1$
 - \vdots

Real equivalence

- two expressions are **real equivalent** if one can be reduced to the other using any of the field identities
 - all of the IEEE identities
 - associativity of addition and multiplication
 - $x(1/x) = 1$
 - \vdots
- real equivalence is **weaker** than IEEE equivalence
 - two real-equivalent programs may produce different results, even when executed on IEEE-compliant platforms
- if arithmetic were infinite-precision, results would be identical
- sometime real equivalence is the best that can be expected
 - suppose program uses MPI_Reduce or MPI_Allreduce
 - reduction operations is floating-point addition

Real equivalence

- two expressions are **real equivalent** if one can be reduced to the other using any of the field identities
 - all of the IEEE identities
 - associativity of addition and multiplication
 - $x(1/x) = 1$
 - \vdots
- real equivalence is **weaker** than IEEE equivalence
 - two real-equivalent programs may produce different results, even when executed on IEEE-compliant platforms
- if arithmetic were infinite-precision, results would be identical
- sometime real equivalence is the best that can be expected
 - suppose program uses MPI_Reduce or MPI_Allreduce
 - reduction operations is floating-point addition
 - **MPI Standard** permits MPI implementation to perform additions in any order
 - order used on one execution could be different than order used on another execution
 - prevents any possibility of IEEE equivalence

Numerical model of diffusion2d

Composite model:

- `diffusion/diffusion_sym.prom`
- `diffusion/diffusion_sym.c`

The path correspondence problem

- the programs may contain branches on expressions that involve the symbolic variables
 - if** ($x_0 \neq 0$) **{...}** **else** **{...}**

Path conditions and domains

- enumerate all paths through the sequential program
 - keeping track of the **path condition** for each path

$$\mathbf{y} = \begin{cases} f_1(\mathbf{x}) & \text{if } p_1(\mathbf{x}) \\ f_2(\mathbf{x}) & \text{if } p_2(\mathbf{x}) \\ \vdots & \vdots \\ f_n(\mathbf{x}) & \text{if } p_n(\mathbf{x}) \end{cases}$$

Path conditions and domains

- enumerate all paths through the sequential program
 - keeping track of the **path condition** for each path

$$\mathbf{y} = \begin{cases} f_1(\mathbf{x}) & \text{if } p_1(\mathbf{x}) \\ f_2(\mathbf{x}) & \text{if } p_2(\mathbf{x}) \\ \vdots & \vdots \\ f_n(\mathbf{x}) & \text{if } p_n(\mathbf{x}) \end{cases}$$

- each p_i determines a **path domain** $D_i = \{\mathbf{x} \mid p_i(\mathbf{x})\}$
- $D_i \cap D_j = \emptyset$ if $i \neq j$
- $\cup_i D_i$ is the whole input space

The method: incorporating path condition

1. construct symbolic model M_{seq} of sequential program
 - input: \mathbf{x} , output: \mathbf{y} , path condition: p
2. construct symbolic model M_{par} of parallel program
 - input: \mathbf{x} , output: \mathbf{y}' , path condition: p
 - using same symbolic table
3. create composite model M :
 - $p \leftarrow \text{true}; M_{\text{seq}}; M_{\text{par}}; \text{assert}(\mathbf{y} = \mathbf{y}')$;
4. use model checker to verify the assertion in M can never be violated

The method: incorporating path condition

1. construct symbolic model M_{seq} of sequential program
 - input: \mathbf{x} , output: \mathbf{y} , path condition: p
2. construct symbolic model M_{par} of parallel program
 - input: \mathbf{x} , output: \mathbf{y}' , path condition: p
 - using same symbolic table
3. create composite model M :
 - $p \leftarrow \text{true}; M_{\text{seq}}; M_{\text{par}}; \text{assert}(\mathbf{y} = \mathbf{y}')$;
4. use model checker to verify the assertion in M can never be violated

The model checker returns either

- Yes: the property holds, or
- No + counterexample:
 - a trace through M_{seq}
 - a trace through M_{par}
 - the values of p , \mathbf{y} , and \mathbf{y}'

Example: Gaussian elimination

- Step 1** Locate the leftmost column of A that does not consist entirely of zeros, if one exists. The top nonzero entry of this column is the pivot.
- Step 2** Interchange the top row with the pivot row, if necessary, so that the entry at the top of the column found in Step 1 is nonzero.
- Step 3** Divide the top row by pivot in order to introduce a leading 1.
- Step 4** Add suitable multiples of the top row to all other rows so that all entries above and below the leading 1 become zero. Repeat.

Gaussian elimination

transforms a matrix to its reduced row-echelon form:

$$\mathbf{x} = \begin{pmatrix} x_0 & x_1 \\ x_2 & x_3 \end{pmatrix} \rightarrow \mathbf{y} = \begin{pmatrix} y_0 & y_1 \\ y_2 & y_3 \end{pmatrix}$$

Gaussian elimination

transforms a matrix to its reduced row-echelon form:

$$\mathbf{x} = \begin{pmatrix} x_0 & x_1 \\ x_2 & x_3 \end{pmatrix} \rightarrow \mathbf{y} = \begin{pmatrix} y_0 & y_1 \\ y_2 & y_3 \end{pmatrix}$$

$$\mathbf{y} = \begin{cases} \begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix} & \text{if } x_0 = 0 \wedge x_2 = 0 \wedge x_1 = 0 \wedge x_3 = 0 \\ \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix} & \text{if } x_0 = 0 \wedge x_2 = 0 \wedge x_1 = 0 \wedge x_3 \neq 0 \\ \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix} & \text{if } x_0 = 0 \wedge x_2 = 0 \wedge x_1 \neq 0 \\ \begin{pmatrix} 1 & x_3/x_2 \\ 0 & 0 \end{pmatrix} & \text{if } x_0 = 0 \wedge x_2 \neq 0 \wedge x_1 = 0 \\ \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} & \text{if } x_0 = 0 \wedge x_2 \neq 0 \wedge x_1 \neq 0 \\ \begin{pmatrix} 1 & x_1/x_0 \\ 0 & 0 \end{pmatrix} & \text{if } x_0 \neq 0 \wedge x_3 - x_2(x_1/x_0) = 0 \\ \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} & \text{if } x_0 \neq 0 \wedge x_3 - x_2(x_1/x_0) \neq 0 \end{cases}$$

Numerical model of Gaussian Elimination

See

- [mpi-spin/examples/gausselim/source/](#)
- [mpi-spin/examples/gausselim/model/](#)